Speed up populational Bayesian inference by combining cross-chain warmup and within-chain parallelization

Yi Zhang[1], William R. Gillespie[1], Ben Bales[2], Aki Vehtari[3]   1. Metrum Research Group, 2. Columbia University, 3. Aalto University

Objectives

- Improve the performance of Bayesian inference of population models by designing a multilevel parallel scheme that combines a cross-chain warmup algorithm in Probabilistic programming language Stan [1] and Torsten's within-chain parallel group ODE solvers [4][6].

- Demonstrate that the new approach significantly improves large PKPD model simulation efficiency.

Cross-chain warmup

The efficacy of standard Stan practice of running a fixed number of warmup iterations is unknown *a priori*. We propose a new algorithm in order to evaluate warmup quality and avoid unnecessary iterations, by checking potential scale reduction coefficients ($\hat{R}$) and effective sample sizes (ESS) [5]:
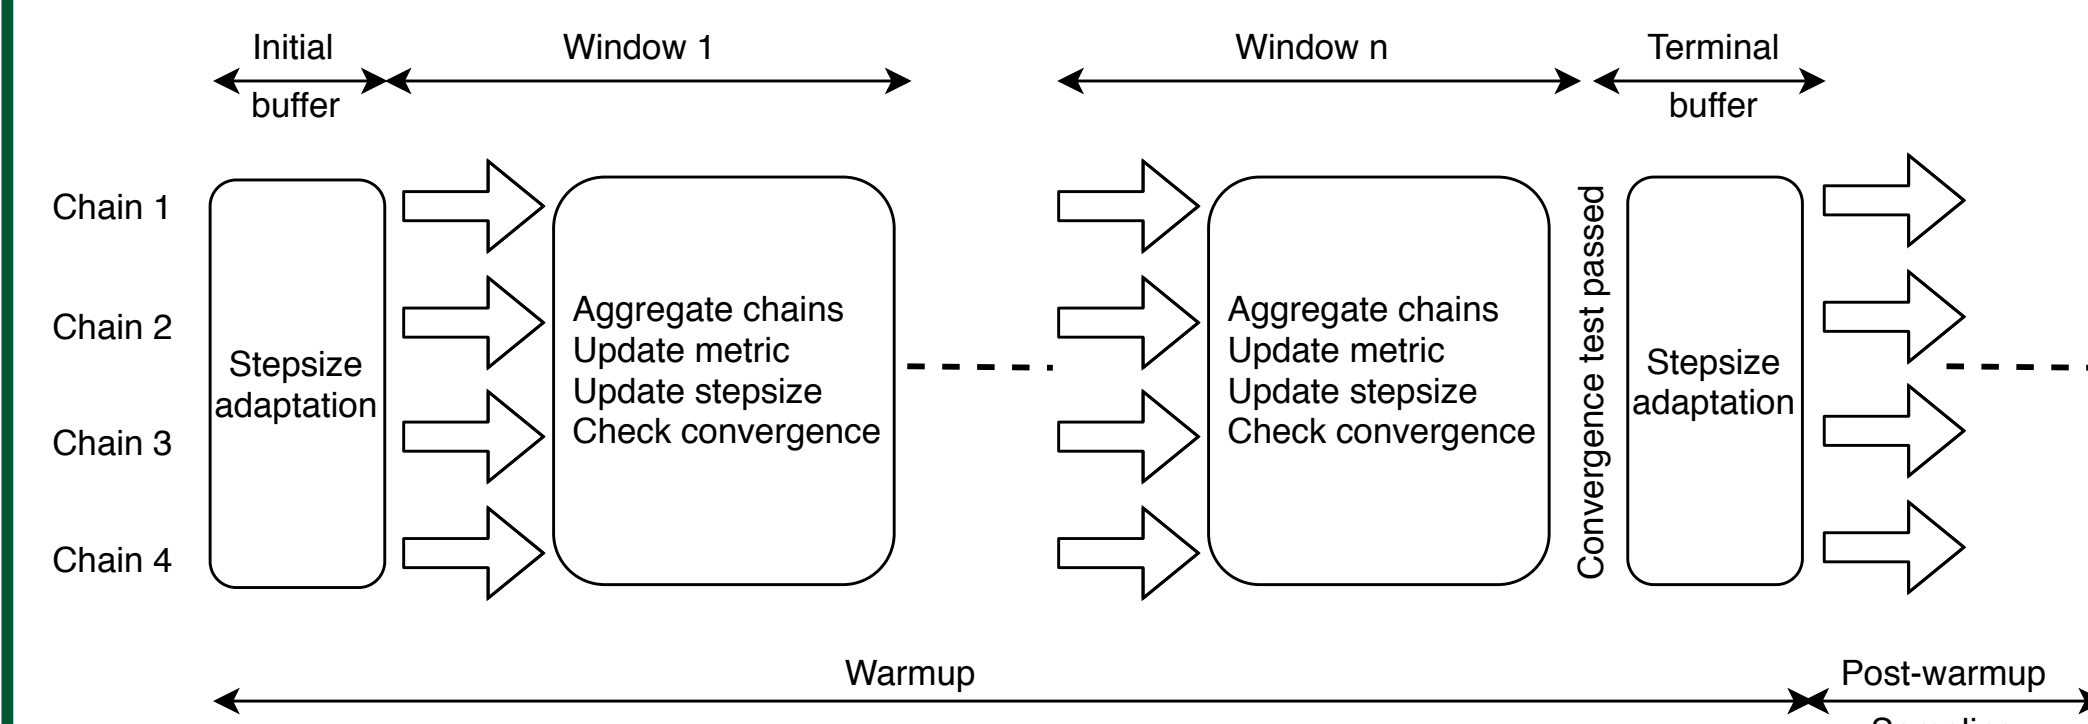

Figure 1: Cross-chain warmup

1. With a fixed window size w, initiate warmup with stepsize adaptation.

2. At the end of a window, aggregate joint posterior probability from all the chains and calculate corresponding $\hat{R}$ and ESS. For example, with default window size $w = 100$, when warmup reaches iteration 200, calculate $\hat{R}^i$ and ESS^i for $i = 1, 2$, so that $\hat{R}^1$ and ESS^1 are based on warmup iteration 1 to 200, and $\hat{R}^2$ and ESS^2 are based on warmup iteration 101 to 200.

3. At the end of window n, with predefined target value $\hat{R}^0$ and ESS^0, from $1, \ldots, n$, select j with maximum ESS^j, and calculate a new metric using samples from corresponding windows. Determine *convergence* by checking if $\hat{R}^j < \hat{R}^0$ and $\text{ESS}^j > \text{ESS}^0$. If converges, move to post-warmup sampling, otherwise repeat step 2.

We call it *cross-chain warmup* as it involves communicating chains.

Benchmarks are performed with different target ESS and regular Stan run(1000 warmup iterations). We run each setup with 10 PRNG seeds and collect average(barplot) and standard deviation(error bar) of the following quantities.

| |
|---|
| total number of leapfrog integration steps in warmup |
| total number of leapfrog integration steps in sampling |
| number of leapfrog integration steps in per each warmup iteration |
| number of leapfrog integration steps in per each sampling iteration |
| minimum ESS_bulk per iteration |
| minimum ESS_tail per iteration |
| minimum ESS_bulk per leapfrog step |
| minimum ESS_tail per leapfrog step |
| maximum wall time(in seconds) |
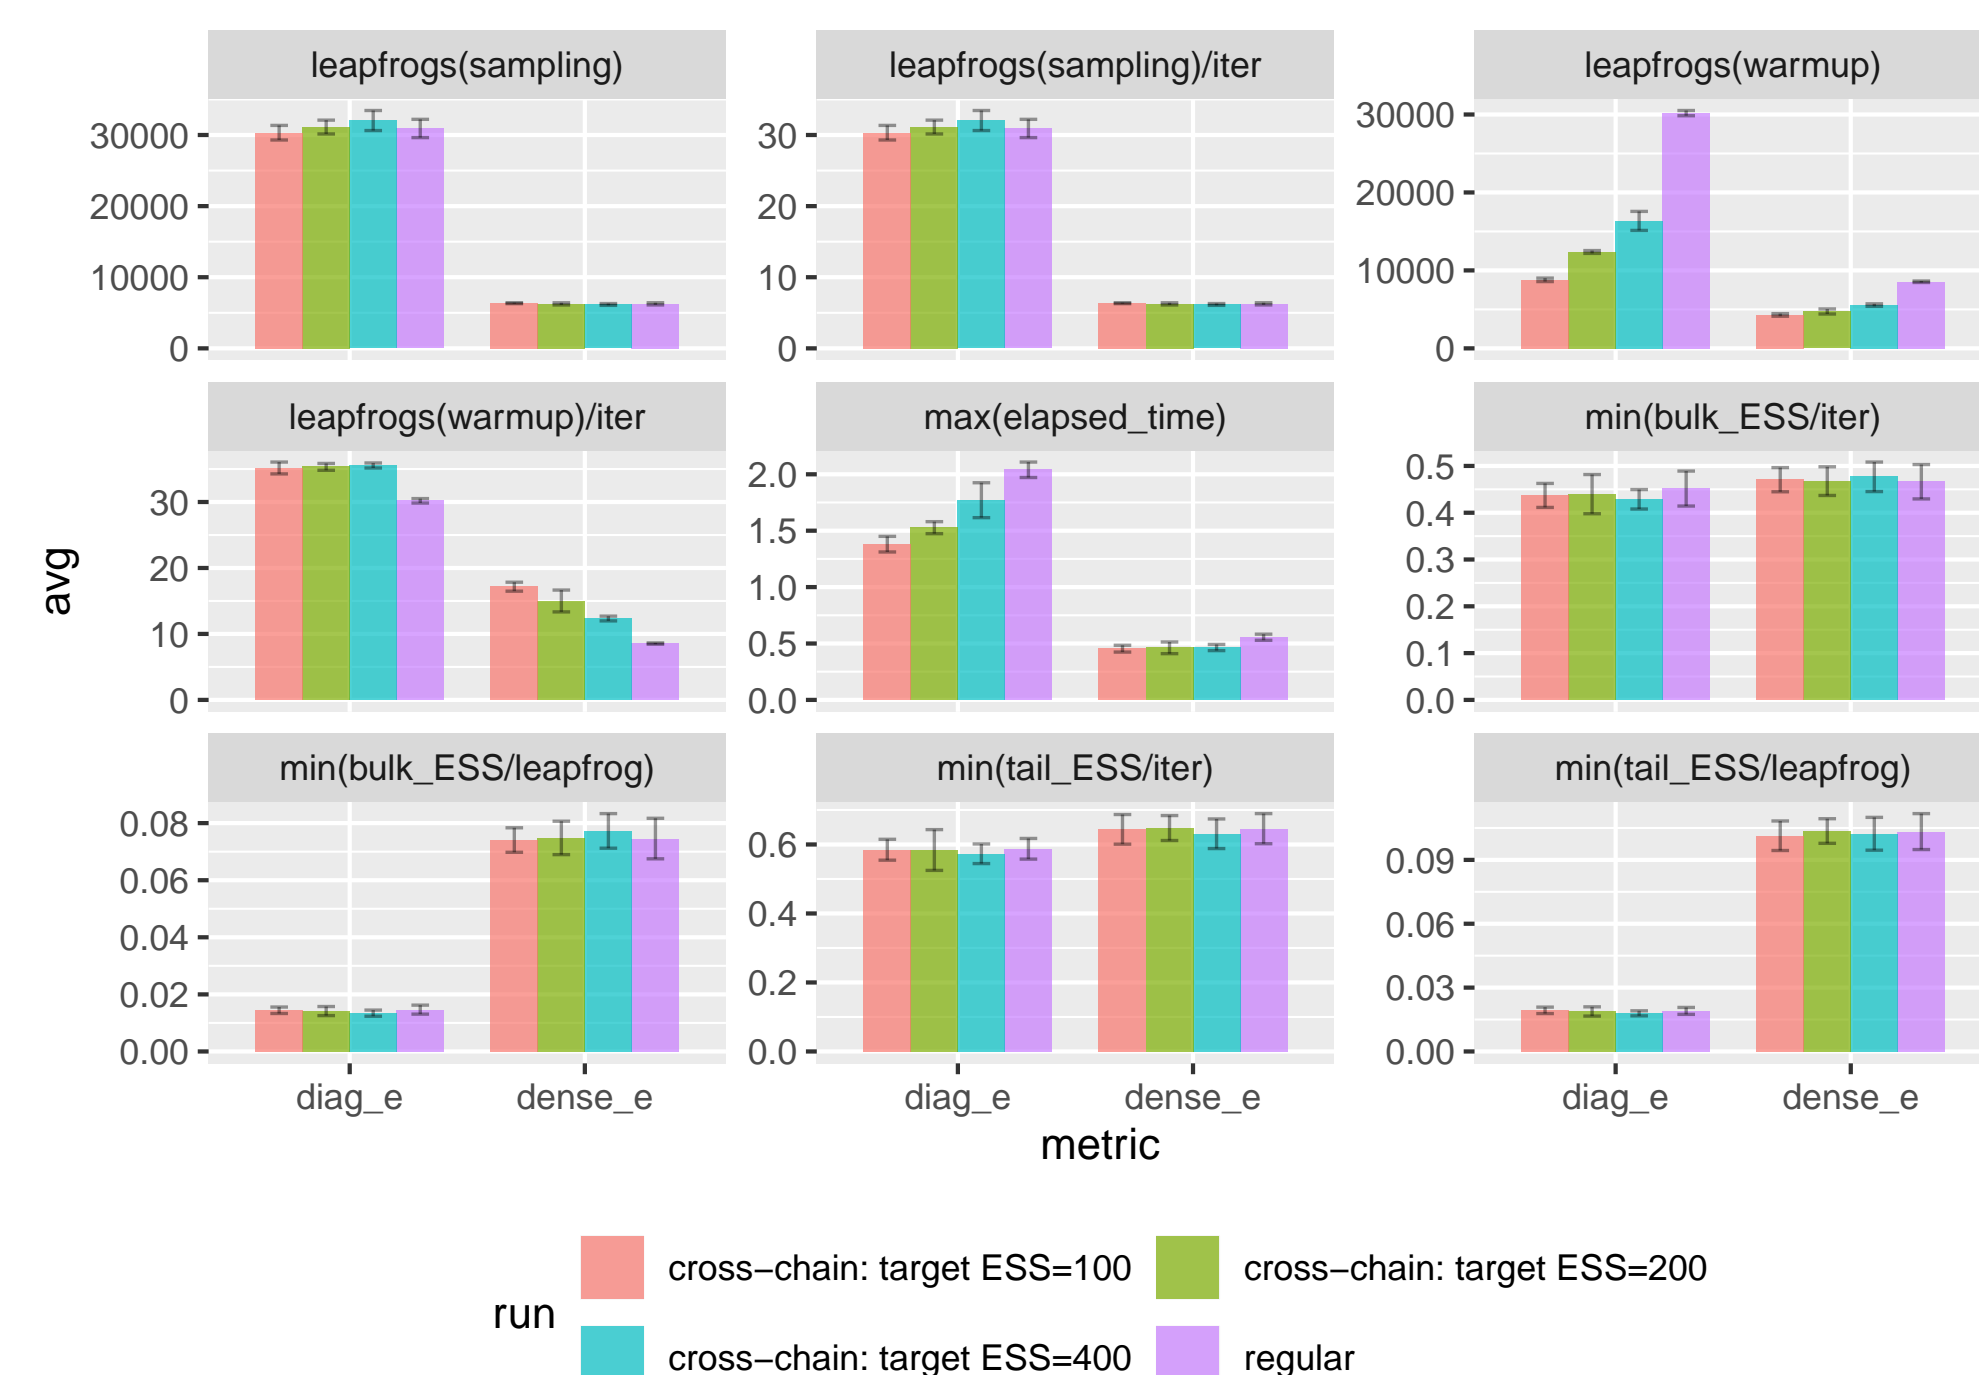
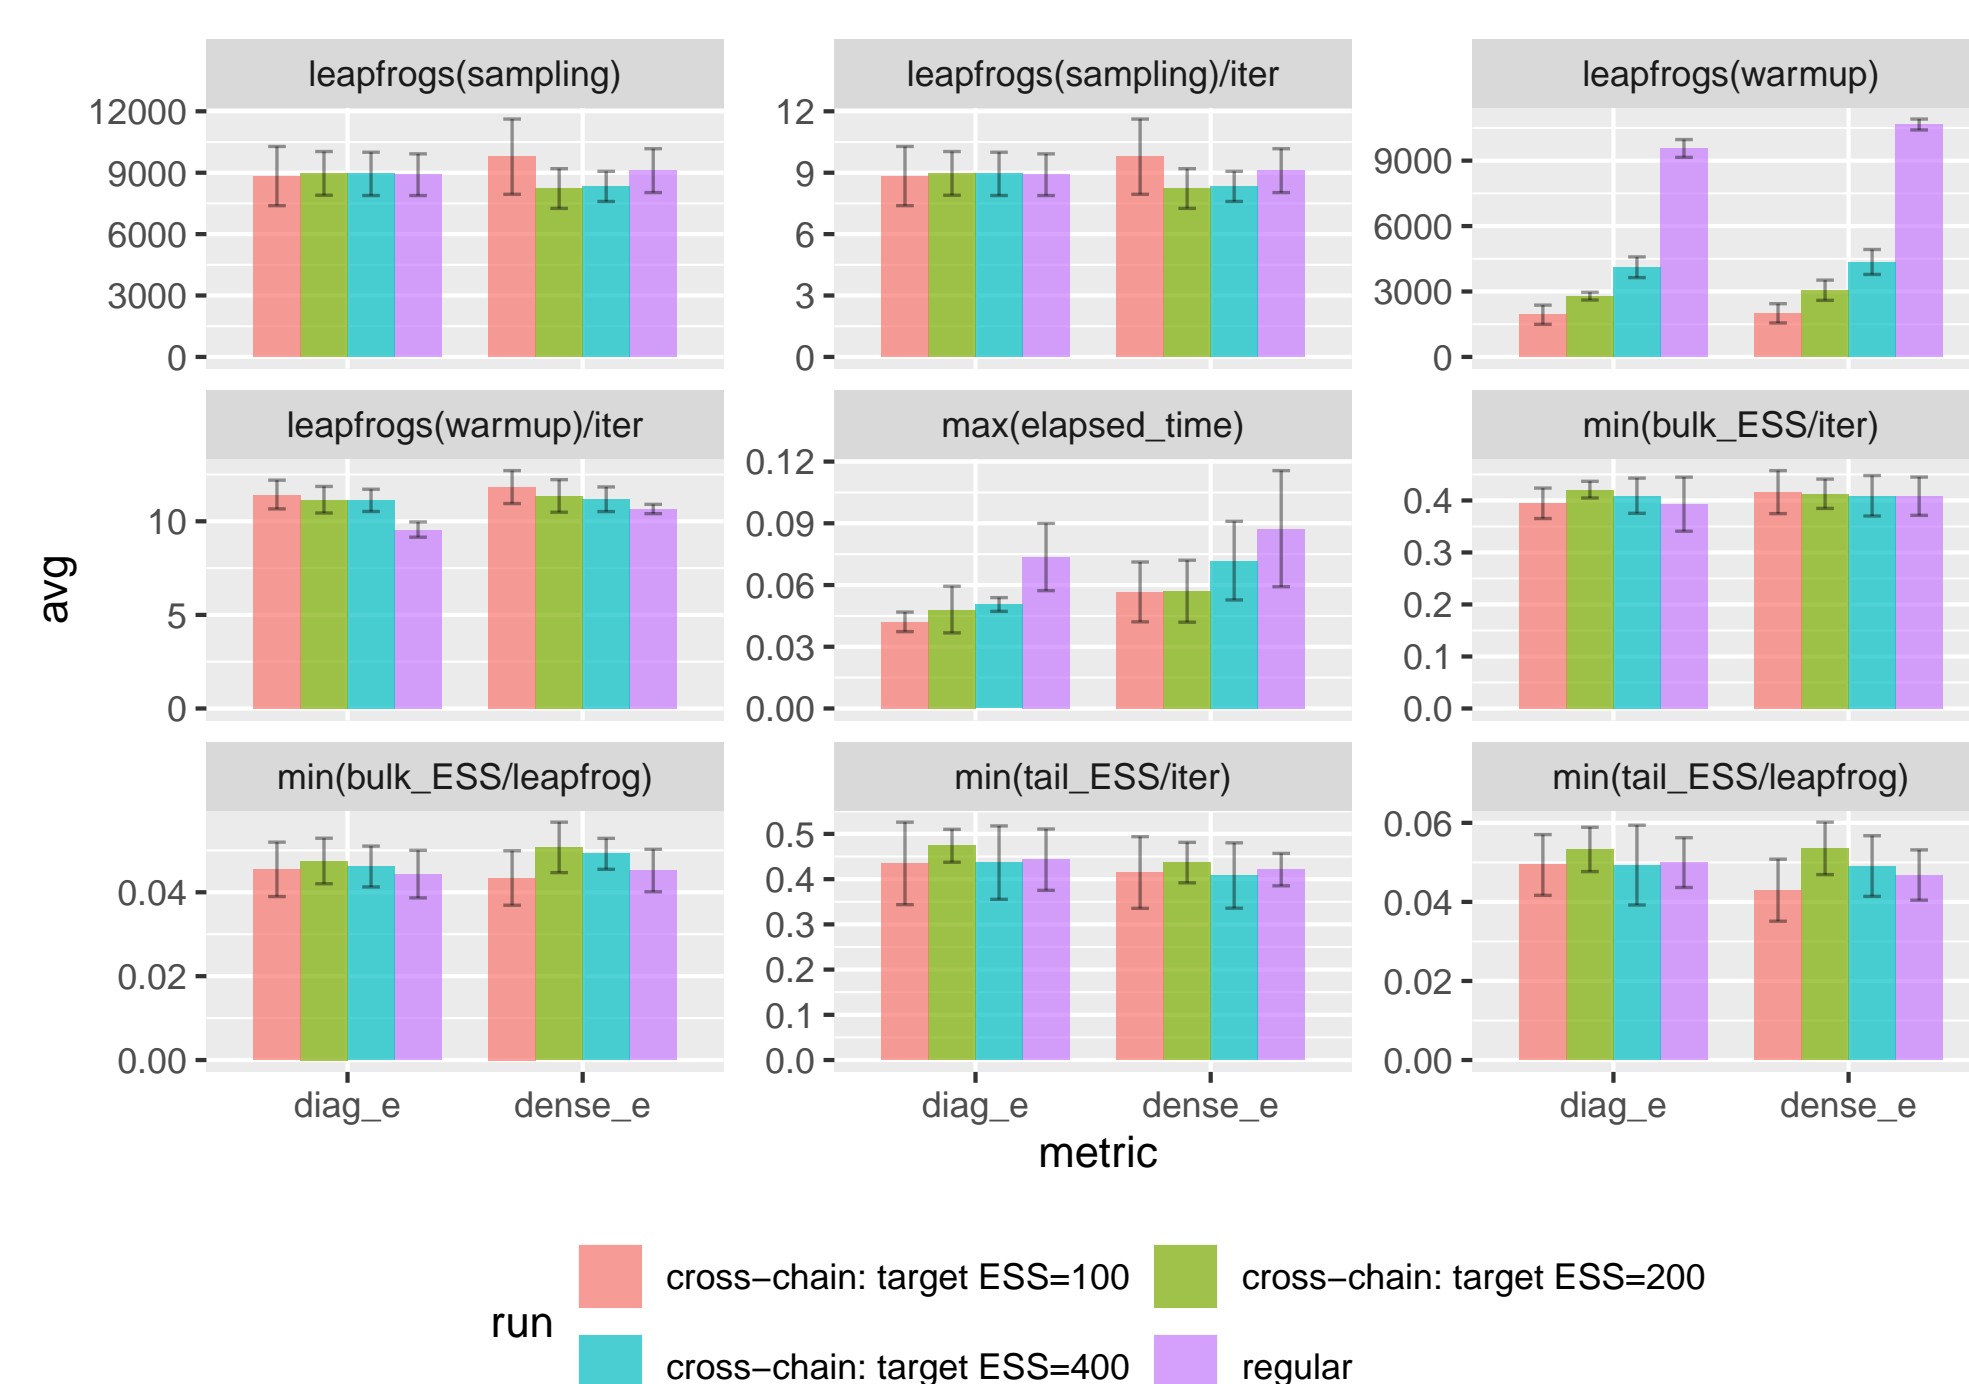**Figure 2:** Cross-chain warmup performance: arK model[3]

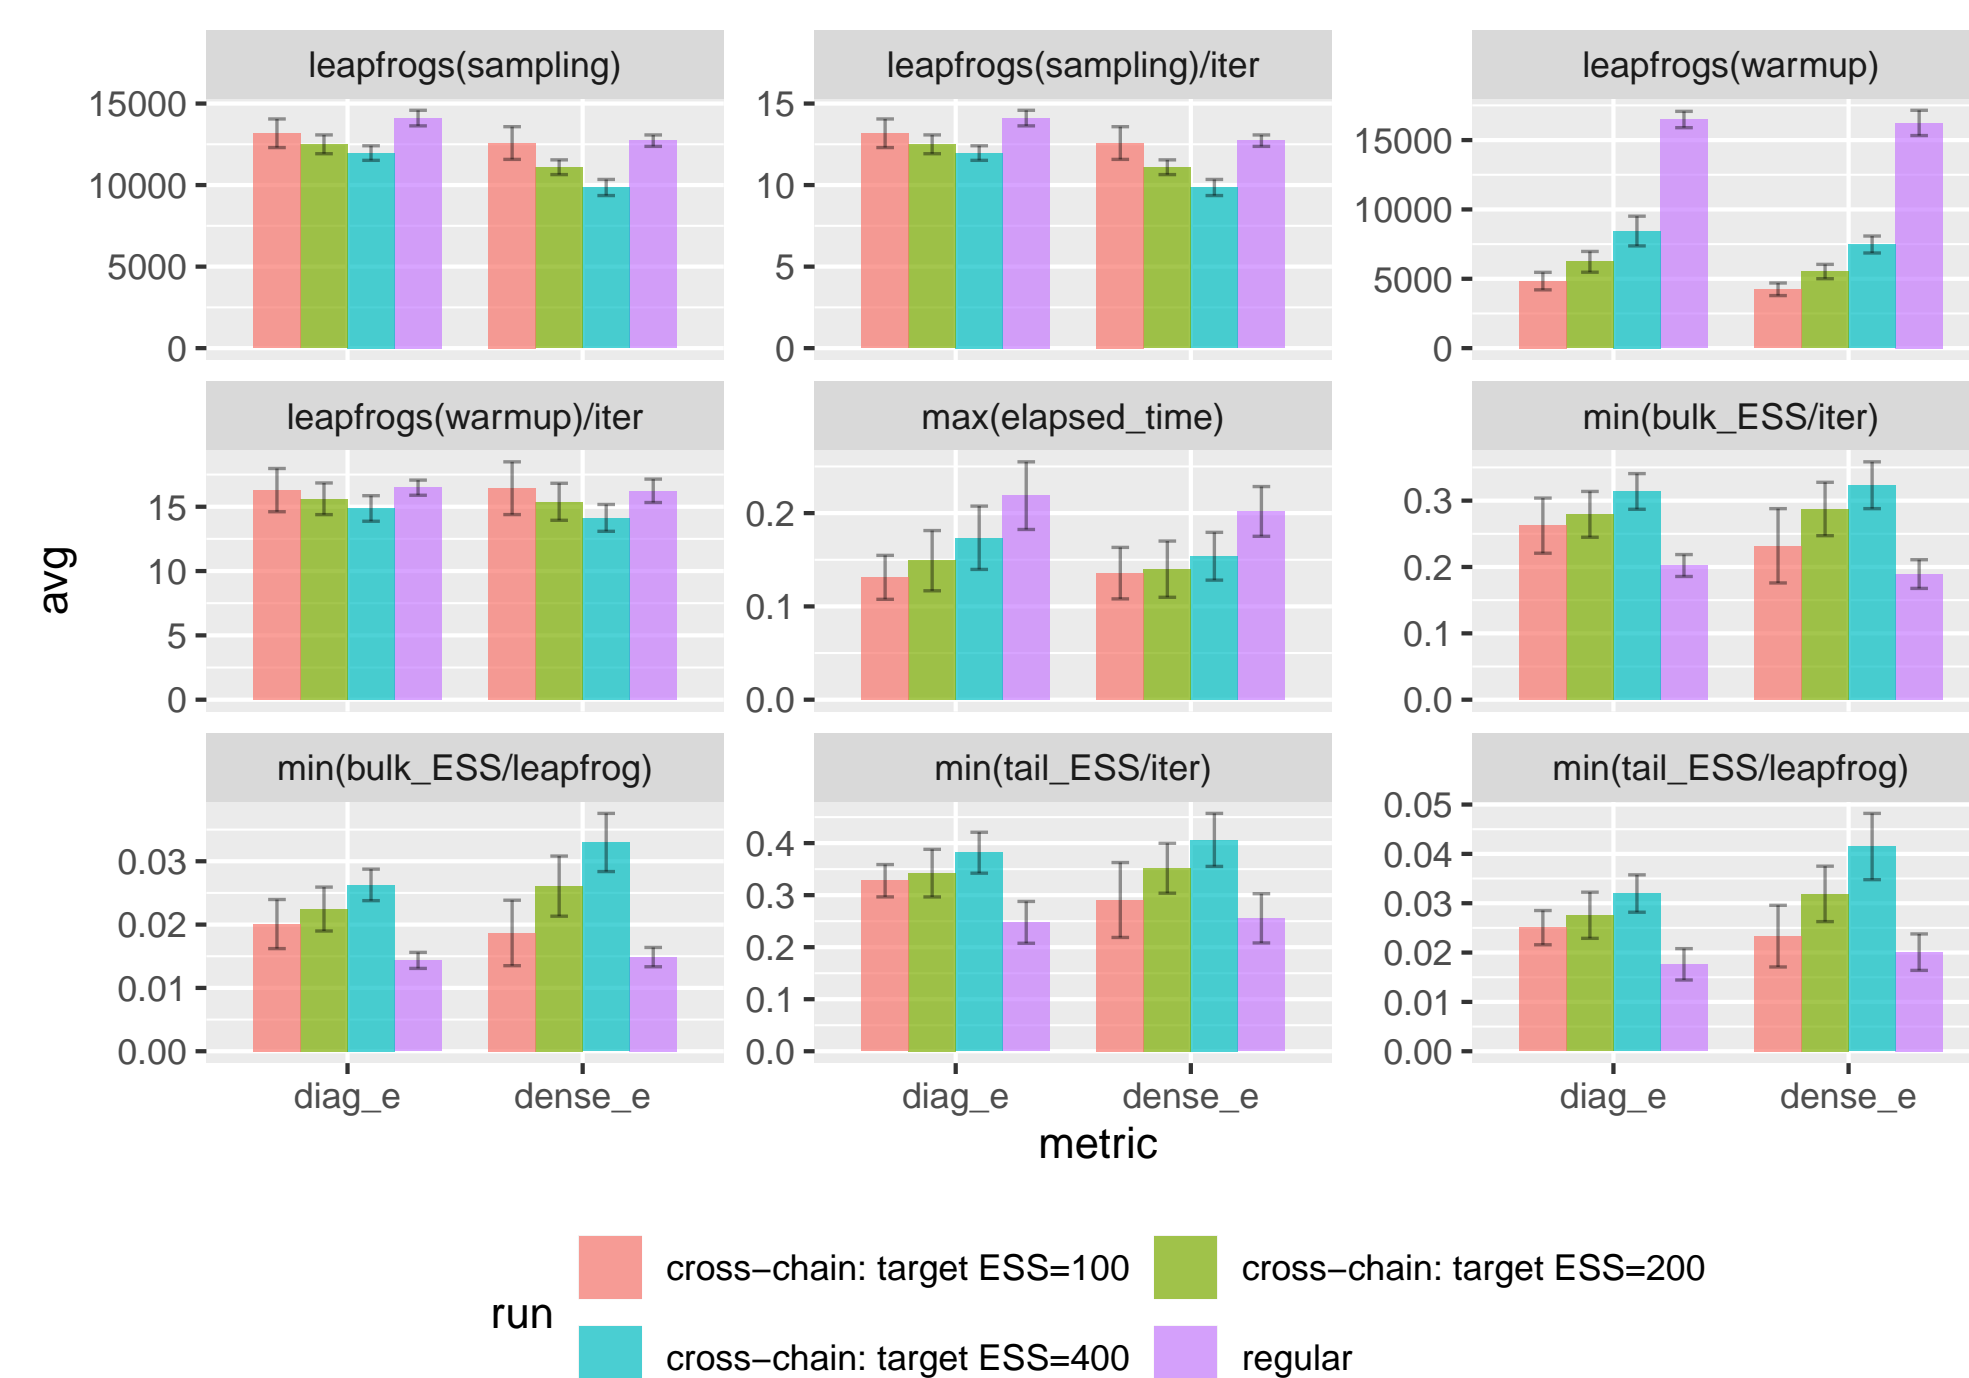**Figure 3:** Cross-chain warmup performance: eight schools model[3]

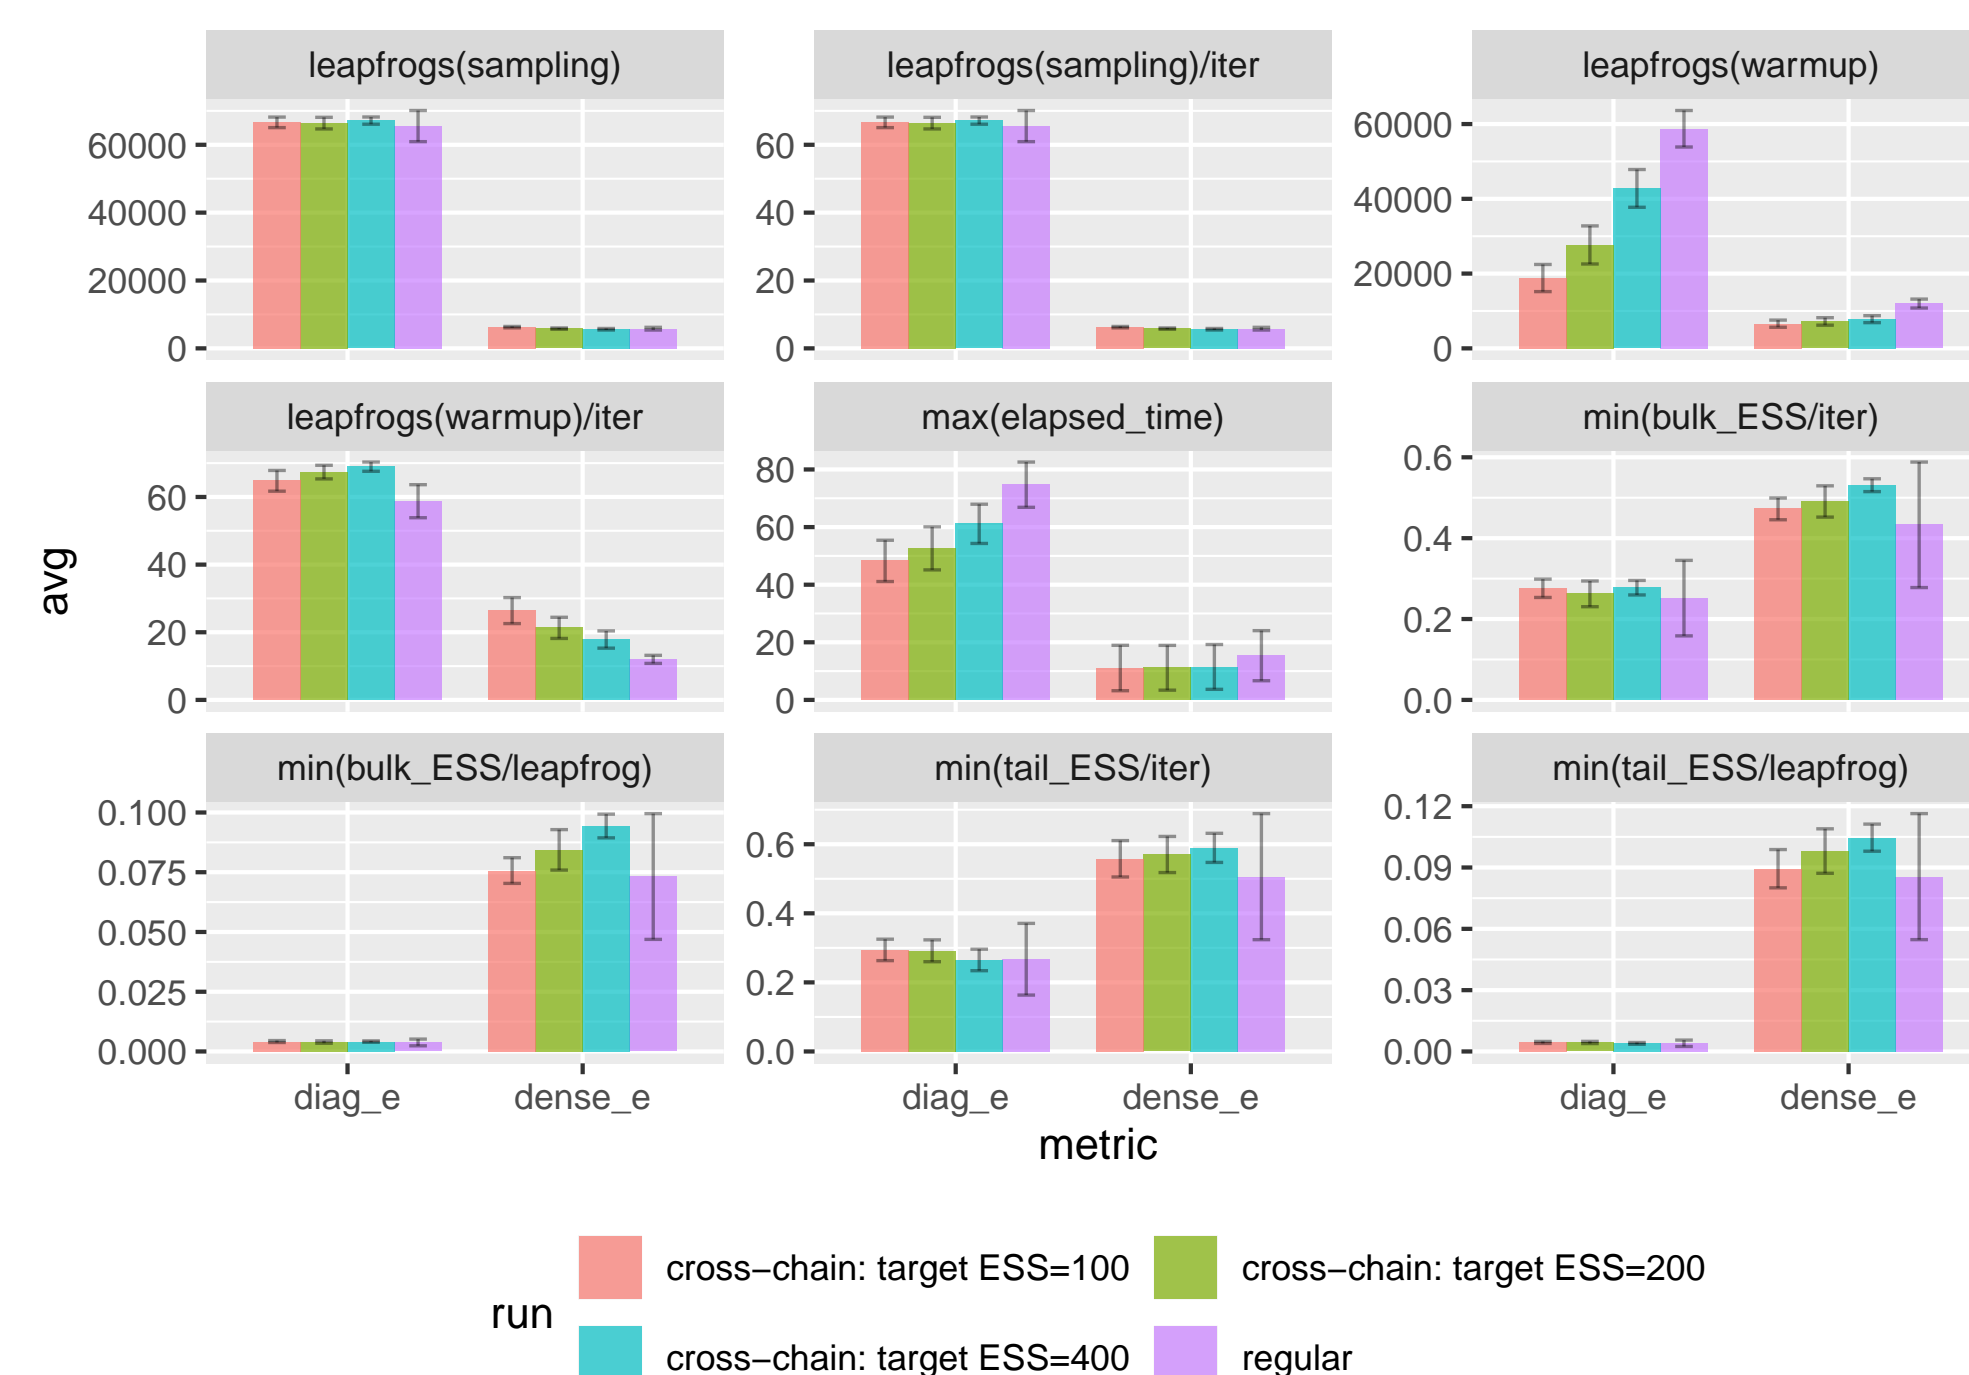**Figure 4:** Cross-chain warmup performance: sblrc-blr model[3]

Figure 5: Cross-chain warmup performance: SIR model[3]

Multilevel parallelization: cross-chain warmup + within-chain parallelization
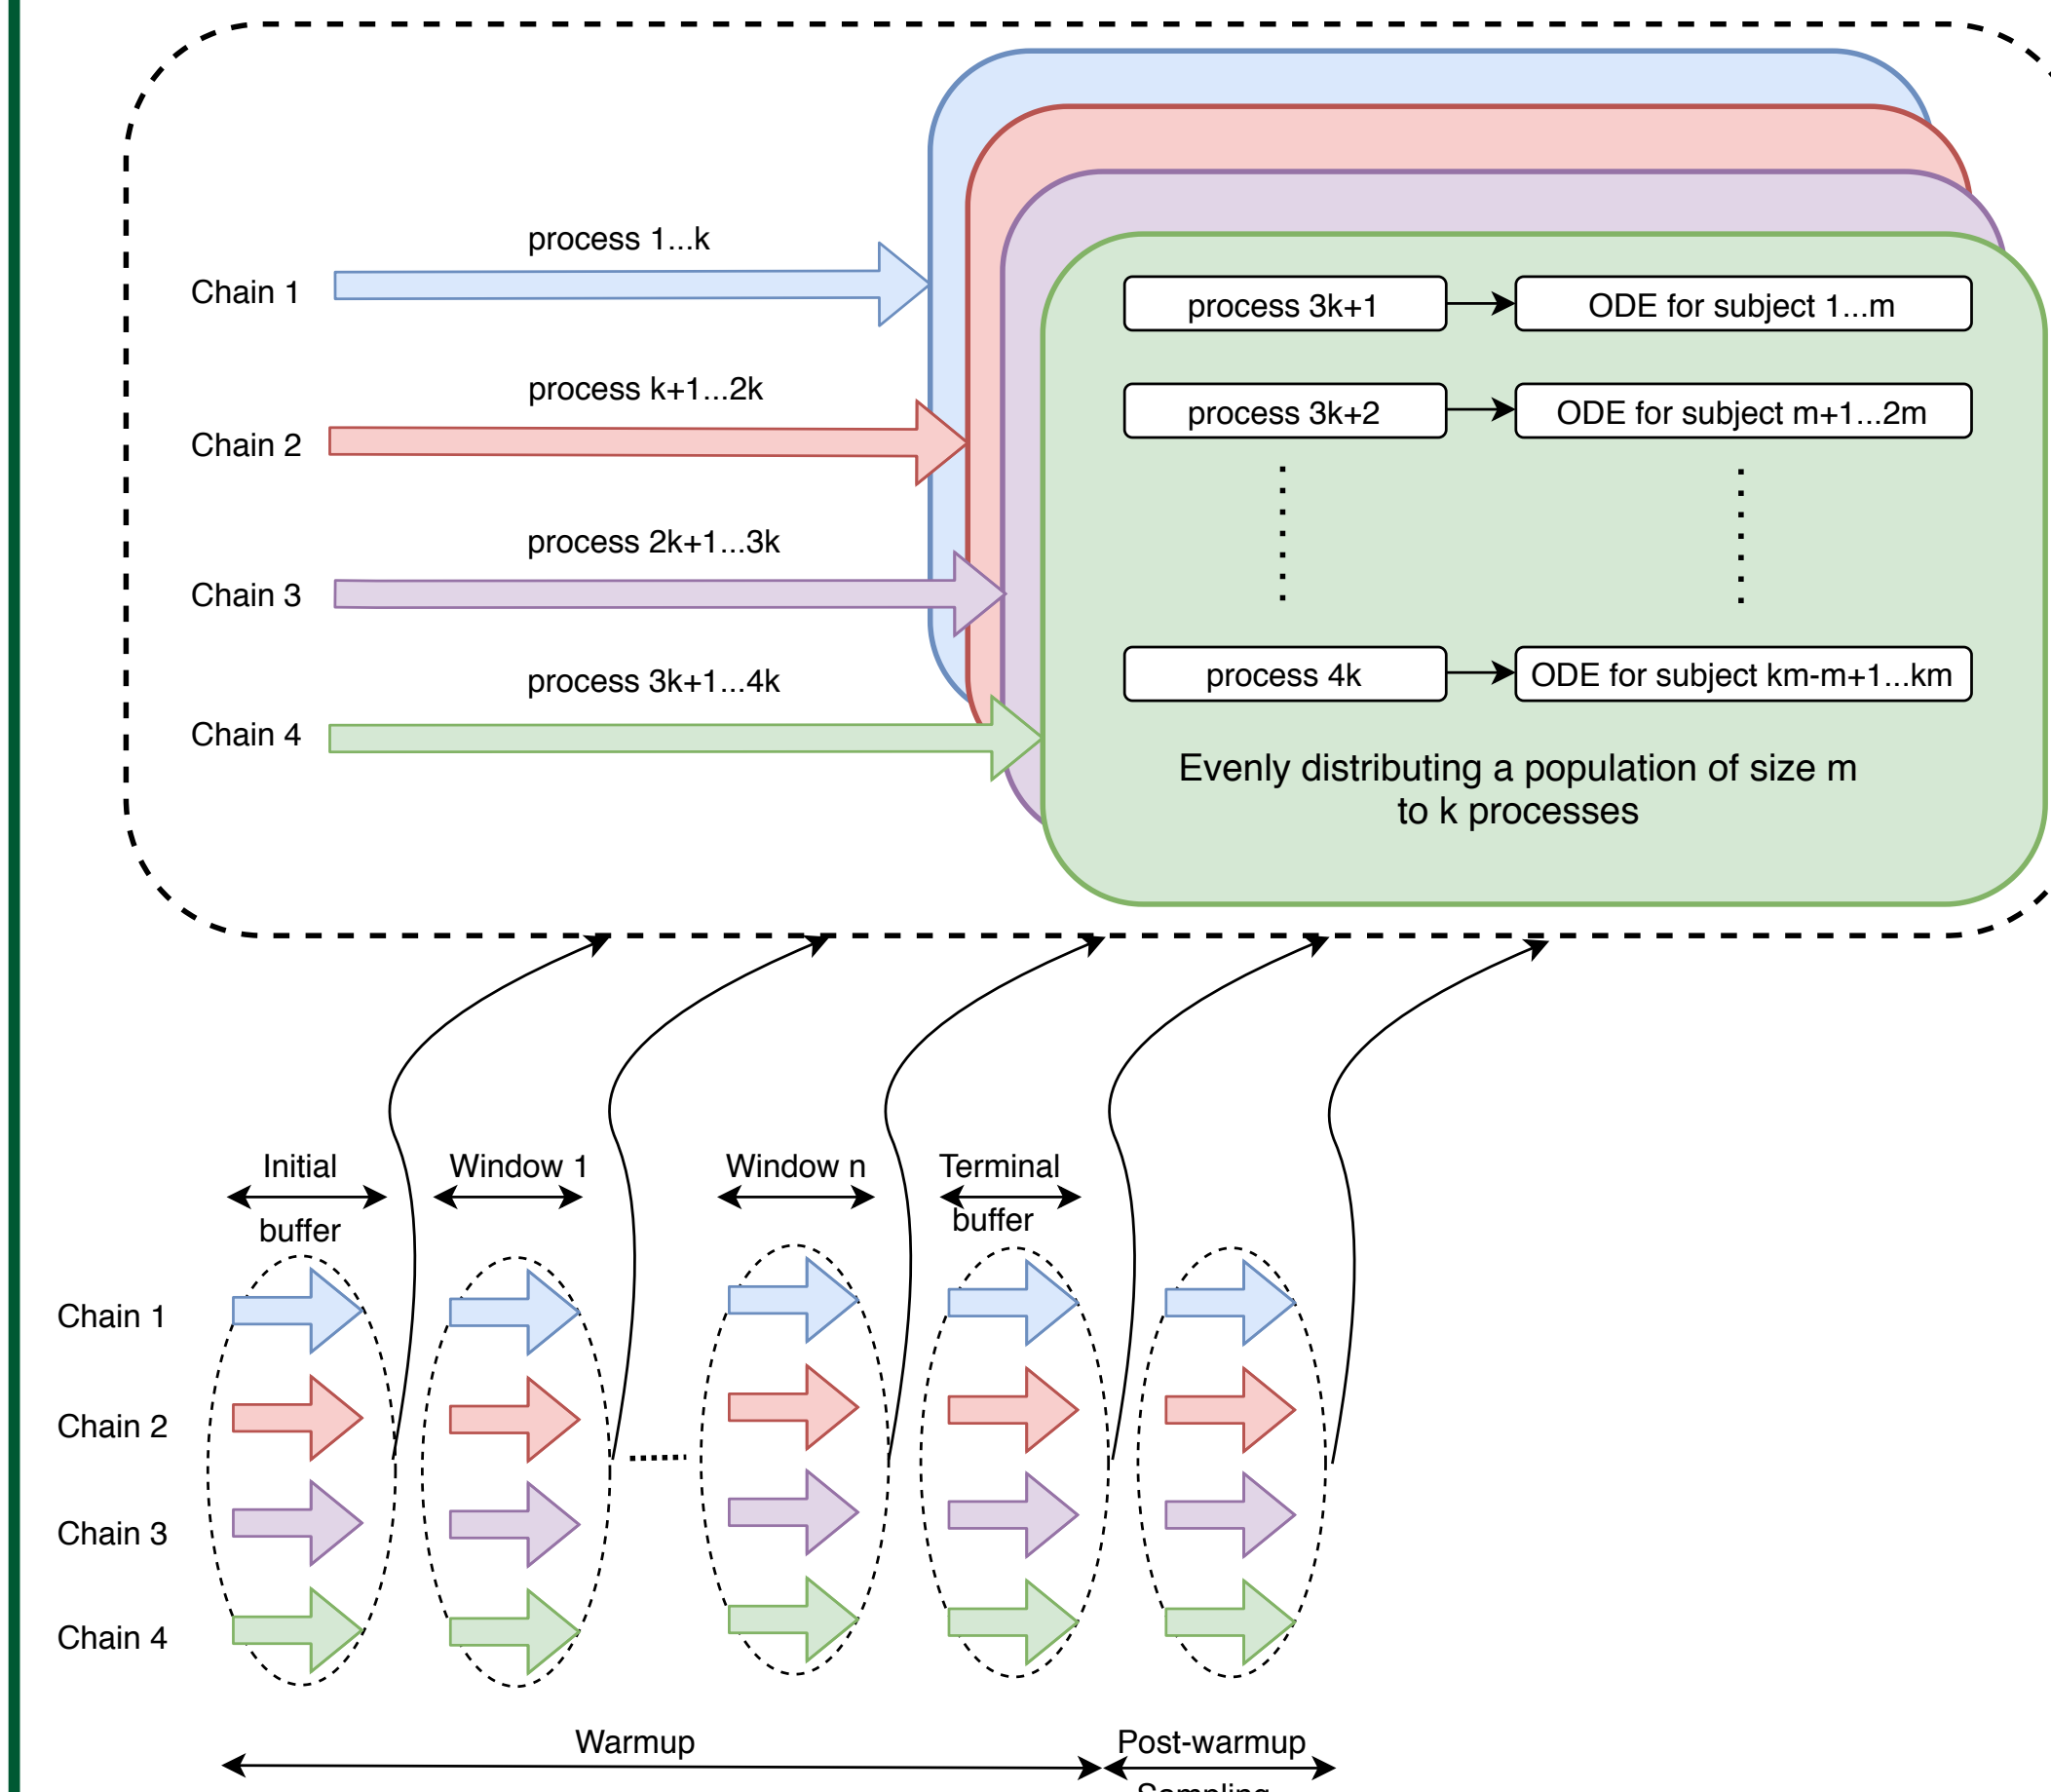
Method

Figure 6: Multilevel parallelism for ODE-based population models. A simplified version of Figure 1, the lower diagram shows the cross-chain warmup through multiple windows. In within-chain parallelization, as shown in the upper diagram, each chain has its own parameter samples(indicated by different colors), and dedicated processes for solving the population model.

Level	Parallel operation	Parallel communication
1	parallel chains with cross-chain warmup	at the end each warmup windo
2	within-chain parallel group ODE solver	when likelihood is evaluated

Table 1: A framework of *multilevel parallelism* for Bayesian inference of population models.

Example

We consider a time-to-event model for the time to the first grade 2+ peripheral neuropathy (PN) event in patients treated with an antibody-drug conjugate (ADC) delivering monomethyl auristatin E (MMAE). We call it Time-To-PN(TTPN) model, and analyze data using a simplified version of the model reported in [2]. We consider three treatment arms: fauxlatuzumab vedotin 1.2, 1.8 and 2.4 mg/kg IV boluses q3w x 6 doses, with 20 patients per treatment arm. In this model, each patient's PK is described by an effective compartment model(one-compartment), and PD by a linear model. The likelihood for time to first 2+ PN event is described by a hazard function that depends on the concentration effect through Weibull distribution. Two unknowns from PK model and the cumulative hazard form a three-component ODE system. Each evaluation of likelihood requires solving this 3-system for every patient.

ODEs corresponding to the entire population are solved by a single call of Torsten function `pmx_solve_group_rk45`. The three parameters of the Torstenn model are:

- k_{e0} in effective compartment model.

- α the coefficient of linear PD model.

- β Weibull distribution scale parameter.

Warmup quality

Table 2 shows cross-chain and regular run performance(target ESS = 400). Consistent with the previous benchmark models, the cross-chain warmup reduce total run time without compromising ESS, leading to 15% wall time improvement.

	Cross-chain	Regular
leapfrogs(warmup)	1.002e+04	1.588e+04
leapfrogs(sampling)	1.709e+04	1.831e+04
leapfrogs(warmup)/iter	1.822e+01	1.588e+01
leapfrogs(sampling)/iter	1.709e+01	1.831e+01
min(bulk_ESS/iter)	2.805e-01	2.340e-01
min(tail_ESS/iter)	3.482e-01	3.205e-01
min(bulk_ESS/leapfrog)	1.641e-02	1.277e-02
min(tail_ESS/leapfrog)	2.037e-02	1.749e-02
max(elapsed_time)	1.702e+03	1.979e+03

Table 2: Cross-chain runs vs regular runs(target ESS=400)
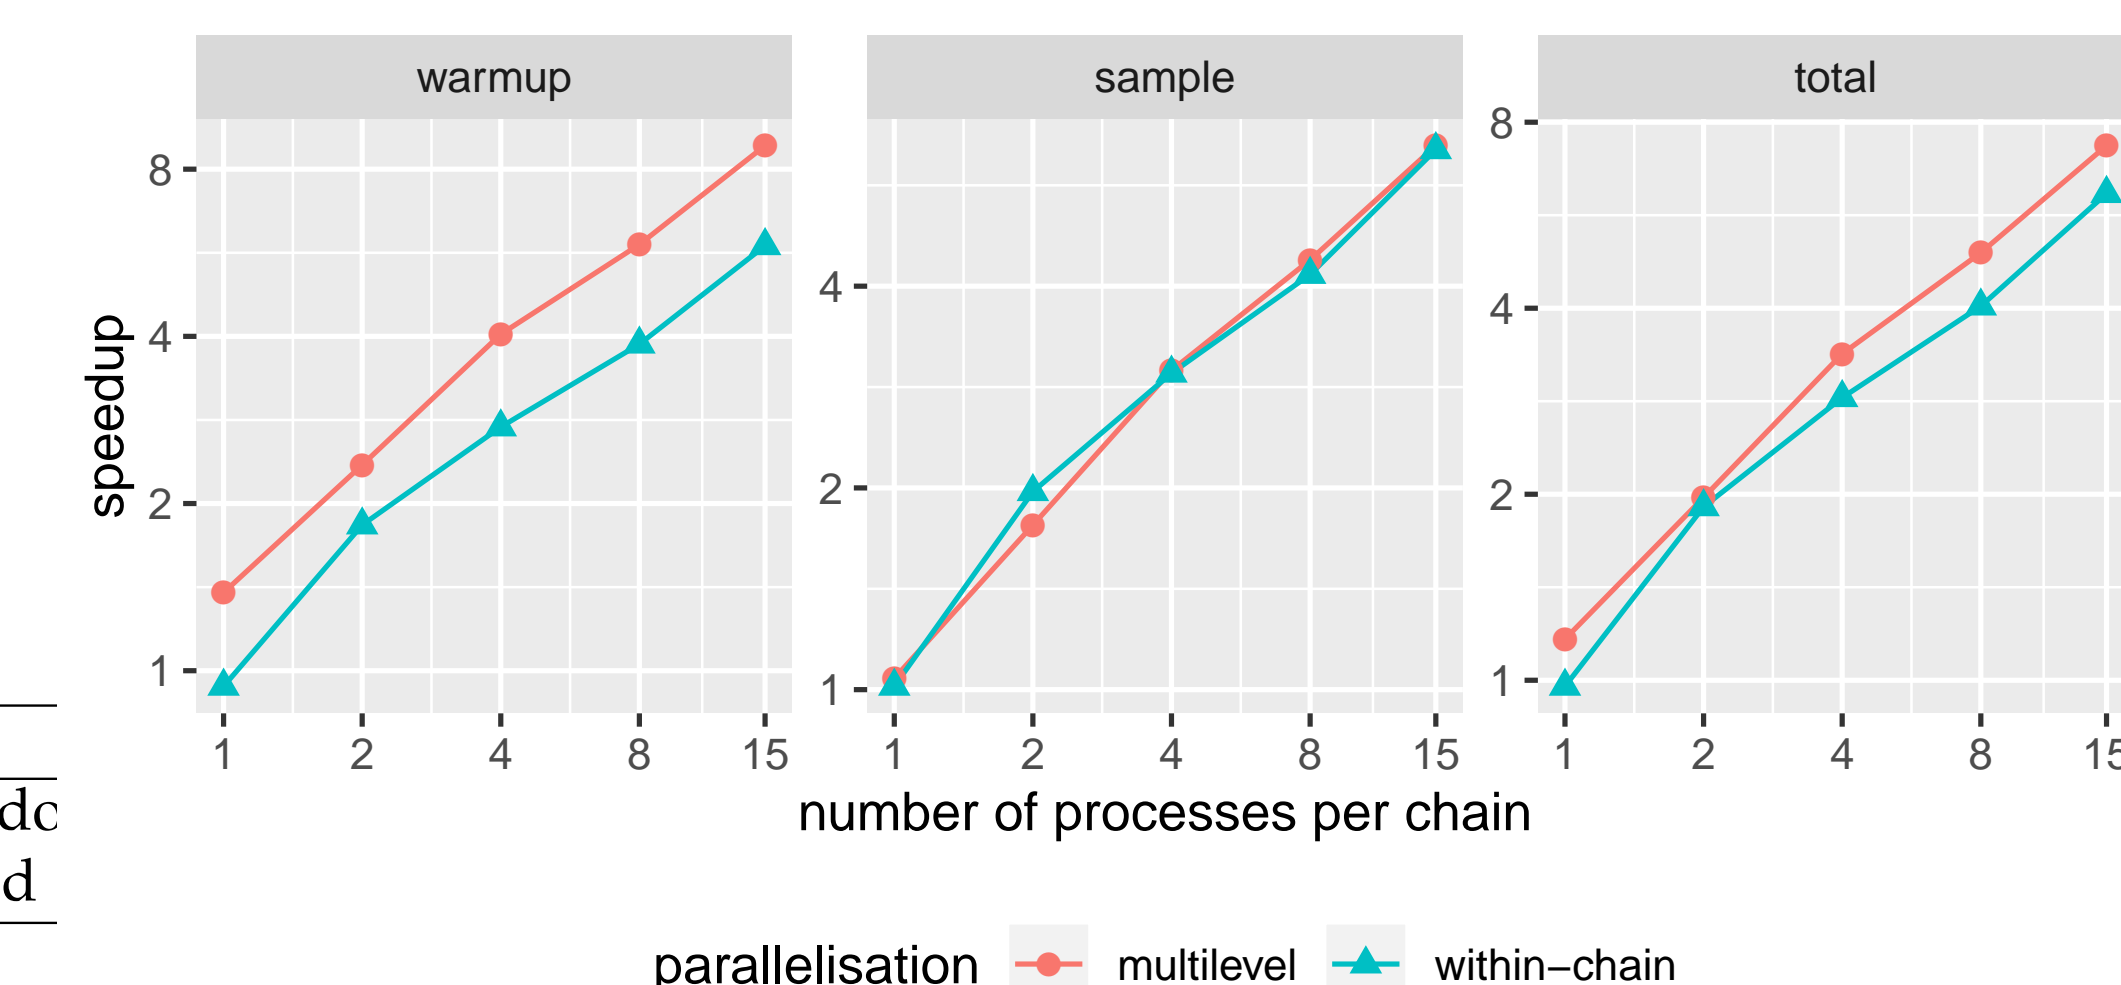
Parallel speedup

Figure 7: Multilevel scheme parallel performance: TTPN model(target ESS=400). Wall time speedup uses regular Stan run as reference. With all runs having 1000 post-warmup sampling iterations, in multilevel runs the number of warmup iterations is determined at runtime, while both within-chain parallel runs and regular Stan runs have 1000 warmup iterations. Among 4 chains in a run, we use the one with maximum total walltime(in seconds) as performance measure, as in practice usually further model evaluation becomes accessible only after all chains finish.

Speedup is investigated by running the model with 4 chains using $n_\text{proc} = 4, 8, 16, 32, 60$ processes. Equivalently, there are $n_\text{proc\_per\_chain} = 1, 2, 4, 8, 15$ processes per chain so that within-chain parallelization can be utilized. With population size 60, each process handles solution of $n_\text{id} = 60, 30, 15, 8, 4$ subjects' ODE system, respectively.

- Both muiltilevel and within-chain-only parallel runs scale near-linearly up to 60 processes(15 processes per chain × 4 chains).

- In the range of $n_\text{proc} = 32, 60, 80$, multilevel runs exhibit a steady ~20% performance improvement, completely contributed by cross-chain warmup.

References

[1] B. Carpenter et al., *Stan: A Probabilistic Programming Language*, Journal of Statistical Software, 76 (2017), pp. 1–32.

[2] D. Lu et al., *Time-to-Event Analysis of Polatuzumab Vedotin-Induced Peripheral Neuropathy to Assist in the Comparison of Clinical Dosing Regimens*, CPT: pharmacometrics & systems pharmacology, 6 (2017), pp. 401–408.

[3] M. Magnusson et al., *A posterior database (PDB) for bayesian inference.* https://github.com/MansMeg/posteriordb.

[4] Torsten Development Team, *Torsten: library of C++ functions that support applications of Stan in Pharmacometrics.* https://github.com/metrumresearchgroup/Torsten.

[5] A. Vehtari et al., *Rank-normalization, folding, and localization: An improved $\hat{R}$ for assessing convergence of MCMC*, arXiv:1903.08008 [stat], (2019). arXiv: 1903.08008.

[6] Y. Zhang and W. R. Gillespie, *Poster: Speed up ode-based modeling using torstens population solvers*, in StanCon 2019, Cambridge, UK, August 2019.

Conclusion and future work

Multilevel parallelism using in Stan and Torsten significantly improves computational efficiency and extends the range of models that may be practically implemented. A natual follow-up of this study is to seek higher efficiency by maintaining target ESS while increasing the number of parallel chains during warmup.

See also

https://github.com/metrumresearchgroup/acop_2020_torsten_parallelization for more benchmark examples and source code.